

Predicting Employee Retention

Objective

The objective of this assignment is to develop a Logistic Regression model. You will be using this model to analyse and predict binary outcomes based on the input data. This assignment aims to enhance understanding of logistic regression, including its assumptions, implementation, and evaluation, to effectively classify and interpret data.

Business Objective

A mid-sized technology company wants to improve its understanding of employee retention to foster a loyal and committed workforce. While the organization has traditionally focused on addressing turnover, it recognises the value of proactively identifying employees likely to stay and understanding the factors contributing to their loyalty.

In this assignment you'll be building a logistic regression model to predict the likelihood of employee retention based on the data such as demographic details, job satisfaction scores, performance metrics, and tenure. The aim is to provide the HR department with actionable insights to strengthen retention strategies, create a supportive work environment, and increase the overall stability and satisfaction of the workforce.

Assignment Tasks

You need to perform the following steps to complete this assignment:

1. Data Understanding
2. Data Cleaning
3. Train Validation Split
4. EDA on training data
5. EDA on validation data [Optional]
6. Feature Engineering
7. Model Building
8. Prediction and Model Evaluation

Data Dictionary

The data has 24 Columns and 74610 Rows. Following data dictionary provides the description for each column present in dataset:

1. Data Understanding

In this step, load the dataset and check basic statistics of the data, including preview of data, dimension of data, column descriptions and data types.

1.0 Import Libraries

```
# Suppress unnecessary warnings
import warnings
warnings.filterwarnings('ignore')

# Import the libraries
import numpy as np
import pandas as pd
```

1.1 Load the Data

```
# Load the dataset
emp_df1 = pd.read_csv("Employee_data.csv")

emp_df1.describe()
```

	Employee ID	Age	Years at Company	Monthly Income \
count	74610.000000	74610.000000	74610.000000	74610.000000
mean	37246.028696	38.529379	15.722638	7344.931417
std	21505.785344	12.082299	11.224059	2596.373589
min	1.000000	18.000000	1.000000	1226.000000
25%	18624.250000	28.000000	7.000000	5652.000000
50%	37239.500000	39.000000	13.000000	7348.500000
75%	55871.750000	49.000000	23.000000	8876.000000
max	74498.000000	59.000000	51.000000	50030.000000

	Number of Promotions	Distance from Home	Number of Dependents
count	74610.000000	72698.000000	74610.000000
mean	0.832958	49.990839	1.657432
std	0.995326	28.519135	1.579862
min	0.000000	1.000000	0.000000
25%	0.000000	25.000000	0.000000
50%	1.000000	50.000000	1.000000
75%	2.000000	75.000000	3.000000
max	4.000000	99.000000	15.000000

	Company Tenure (In Months)
count	72197.000000
mean	55.711899
std	25.392325
min	2.000000

25%	36.000000
50%	56.000000
75%	76.000000
max	128.000000

Check the first few entries

emp_df1.head()

	Employee ID	Age	Gender	Years at Company	Job Role	Monthly Income \
0	8410	31	Male	19	Education	5390
1	64756	59	Female	4	Media	5534
2	30257	24	Female	10	Healthcare	8159
3	65791	36	Female	7	Education	3989
4	65026	56	Male	41	Education	4821

	Work-Life Balance	Job Satisfaction	Performance Rating	Number of Promotions \
0	Excellent	Medium	Average	2
1	Poor	High	Low	3
2	Good	High	Low	0
3	Good	High	High	1
4	Fair	Very High	Average	0

	... Number of Dependents	Job Level	Company Size \
0	...	0 Mid	Medium
1	...	3 Mid	Medium
2	...	3 Mid	Medium
3	...	2 Mid	Small
4	...	0 Senior	Medium

	Company Tenure (In Months)	Remote Work	Leadership Opportunities \
0	89.0	No	No
1	21.0	No	No
2	74.0	No	No
3	50.0	Yes	No
4	68.0	No	No

	Innovation Opportunities	Company Reputation	Employee Recognition	Attrition
--	--------------------------	--------------------	----------------------	-----------

0	No	Excellent	Medium
Stayed			
1	No	Fair	Low
Stayed			
2	No	Poor	Low
Stayed			
3	No	Good	Medium
Stayed			
4	No	Fair	Medium
Stayed			

[5 rows x 24 columns]

Inspect the shape of the dataset

emp_df1.shape

(74610, 24)

Inspect the different columns in the dataset

emp_df1.columns

```
Index(['Employee ID', 'Age', 'Gender', 'Years at Company', 'Job Role',
      'Monthly Income', 'Work-Life Balance', 'Job Satisfaction',
      'Performance Rating', 'Number of Promotions', 'Overtime',
      'Distance from Home', 'Education Level', 'Marital Status',
      'Number of Dependents', 'Job Level', 'Company Size',
      'Company Tenure (In Months)', 'Remote Work', 'Leadership
      Opportunities',
      'Innovation Opportunities', 'Company Reputation',
      'Employee Recognition', 'Attrition'],
      dtype='object')
```

1.2 Check the basic statistics

emp_df1.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 74610 entries, 0 to 74609

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	Employee ID	74610 non-null	int64
1	Age	74610 non-null	int64
2	Gender	74610 non-null	object
3	Years at Company	74610 non-null	int64
4	Job Role	74610 non-null	object
5	Monthly Income	74610 non-null	int64
6	Work-Life Balance	74610 non-null	object
7	Job Satisfaction	74610 non-null	object
8	Performance Rating	74610 non-null	object

9	Number of Promotions	74610	non-null	int64
10	Overtime	74610	non-null	object
11	Distance from Home	72698	non-null	float64
12	Education Level	74610	non-null	object
13	Marital Status	74610	non-null	object
14	Number of Dependents	74610	non-null	int64
15	Job Level	74610	non-null	object
16	Company Size	74610	non-null	object
17	Company Tenure (In Months)	72197	non-null	float64
18	Remote Work	74610	non-null	object
19	Leadership Opportunities	74610	non-null	object
20	Innovation Opportunities	74610	non-null	object
21	Company Reputation	74610	non-null	object
22	Employee Recognition	74610	non-null	object
23	Attrition	74610	non-null	object

dtypes: float64(2), int64(6), object(16)
memory usage: 13.7+ MB

1.3 Check the data type of columns

Check the info to see the types of the feature variables and the null values present
emp_df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74610 entries, 0 to 74609
Data columns (total 24 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Employee ID                               74610 non-null  int64
1   Age                                         74610 non-null  int64
2   Gender                                     74610 non-null  object
3   Years at Company                           74610 non-null  int64
4   Job Role                                   74610 non-null  object
5   Monthly Income                             74610 non-null  int64
6   Work-Life Balance                           74610 non-null  object
7   Job Satisfaction                           74610 non-null  object
8   Performance Rating                         74610 non-null  object
9   Number of Promotions                       74610 non-null  int64
10  Overtime                                    74610 non-null  object
11  Distance from Home                         72698 non-null  float64
12  Education Level                             74610 non-null  object
13  Marital Status                             74610 non-null  object
14  Number of Dependents                       74610 non-null  int64
15  Job Level                                   74610 non-null  object
16  Company Size                               74610 non-null  object
17  Company Tenure (In Months)                 72197 non-null  float64
18  Remote Work                                74610 non-null  object
19  Leadership Opportunities                   74610 non-null  object
20  Innovation Opportunities                   74610 non-null  object
```

```
21 Company Reputation      74610 non-null object
22 Employee Recognition    74610 non-null object
23 Attrition               74610 non-null object
dtypes: float64(2), int64(6), object(16)
memory usage: 13.7+ MB
```

2. Data Cleaning [15 marks]

2.1 Handle the missing values [10 marks]

2.1.1 Check the number of missing values [2 Mark]

```
# Check the number of missing values in each column
emp_df1.isnull().sum()
```

```
Employee ID      0
Age              0
Gender           0
Years at Company 0
Job Role         0
Monthly Income   0
Work-Life Balance 0
Job Satisfaction 0
Performance Rating 0
Number of Promotions 0
Overtime         0
Distance from Home 1912
Education Level  0
Marital Status   0
Number of Dependents 0
Job Level        0
Company Size     0
Company Tenure (In Months) 2413
Remote Work      0
Leadership Opportunities 0
Innovation Opportunities 0
Company Reputation 0
Employee Recognition 0
Attrition        0
dtype: int64
```

2.1.2 Check the percentage of missing values [2 Marks]

```
# Check the percentage of missing values in each column
missing_percentage = (emp_df1.isnull().sum() / len(emp_df1)) * 100

missing_percentage
```

Employee ID	0.000000
Age	0.000000
Gender	0.000000
Years at Company	0.000000
Job Role	0.000000
Monthly Income	0.000000
Work-Life Balance	0.000000
Job Satisfaction	0.000000
Performance Rating	0.000000
Number of Promotions	0.000000
Overtime	0.000000
Distance from Home	2.562659
Education Level	0.000000
Marital Status	0.000000
Number of Dependents	0.000000
Job Level	0.000000
Company Size	0.000000
Company Tenure (In Months)	3.234151
Remote Work	0.000000
Leadership Opportunities	0.000000
Innovation Opportunities	0.000000
Company Reputation	0.000000
Employee Recognition	0.000000
Attrition	0.000000
dtype: float64	

2.1.3 Handle rows with missing values [4 Marks]

```
# Handle the missing value rows in the column
emp_df2 = emp_df1.dropna()

emp_df2.shape

(70635, 24)
```

2.1.4 Check percentage of remaining data after missing values are removed [2 Mark]

```
remaining_emp_percentage = (len(emp_df2) / len(emp_df1)) * 100

remaining_emp_percentage

94.67229593888219
```

2.2 Identify and handle redundant values within categorical columns (if any) [3 marks]

Examine the categorical columns to determine if any value or column needs to be treated

```

# Write a function to display the categorical columns with their
unique values and check for redundant values
def display_categorical_columns(emp_df):
    categorical_columns =
emp_df.select_dtypes(include=['object']).columns # Filter categorical
columns
    for column in categorical_columns:
        unique_values = emp_df[column].unique() # Get unique values
        print(f"Column: {column}")
        print(f"Unique Values: {unique_values}")
        print(f"Number of Unique Values: {len(unique_values)}")
        print("-" * 50)

# Check the data
display_categorical_columns(emp_df2)

```

```

Column: Gender
Unique Values: ['Male' 'Female']
Number of Unique Values: 2
-----
Column: Job Role
Unique Values: ['Education' 'Media' 'Healthcare' 'Technology'
'Finance']
Number of Unique Values: 5
-----
Column: Work-Life Balance
Unique Values: ['Excellent' 'Poor' 'Good' 'Fair']
Number of Unique Values: 4
-----
Column: Job Satisfaction
Unique Values: ['Medium' 'High' 'Very High' 'Low']
Number of Unique Values: 4
-----
Column: Performance Rating
Unique Values: ['Average' 'Low' 'High' 'Below Average']
Number of Unique Values: 4
-----
Column: Overtime
Unique Values: ['No' 'Yes']
Number of Unique Values: 2
-----
Column: Education Level
Unique Values: ['Associate Degree' 'Masterâ€™s Degree' 'Bachelorâ€™s
Degree'
'High School' 'PhD']
Number of Unique Values: 5
-----
Column: Marital Status
Unique Values: ['Married' 'Divorced' 'Single']
Number of Unique Values: 3

```



```

-----
Column: Job Level
Unique Values: ['Mid' 'Senior' 'Entry']
Number of Unique Values: 3
-----
Column: Company Size
Unique Values: ['Medium' 'Small' 'Large']
Number of Unique Values: 3
-----
Column: Remote Work
Unique Values: ['No' 'Yes']
Number of Unique Values: 2
-----
Column: Leadership Opportunities
Unique Values: ['No' 'Yes']
Number of Unique Values: 2
-----
Column: Innovation Opportunities
Unique Values: ['No' 'Yes']
Number of Unique Values: 2
-----
Column: Company Reputation
Unique Values: ['Excellent' 'Fair' 'Poor' 'Good']
Number of Unique Values: 4
-----
Column: Employee Recognition
Unique Values: ['Medium' 'Low' 'High' 'Very High']
Number of Unique Values: 4
-----
Column: Attrition
Unique Values: ['Stayed' 'Left']
Number of Unique Values: 2
-----

```

2.3 Drop redundant columns [2 marks]

```

# Drop redundant columns which are not required for modelling
columns_to_drop = ['Employee ID', 'Company Tenure (In Months)']
emp_df3 = emp_df2.drop(columns=columns_to_drop, axis=1)

# Check first few rows of data
emp_df3.head()

```

	Age	Gender	Years at Company	Job Role	Monthly Income \
0	31	Male	19	Education	5390
1	59	Female	4	Media	5534
2	24	Female	10	Healthcare	8159
3	36	Female	7	Education	3989
4	56	Male	41	Education	4821

Work-Life Balance	Job Satisfaction	Performance Rating	Number of Promotions \
0	Excellent	Medium	Average
2			
1	Poor	High	Low
3			
2	Good	High	Low
0			
3	Good	High	High
1			
4	Fair	Very High	Average
0			

Overtime ...	Marital Status	Number of Dependents	Job Level
Company Size \			
0	No ...	Married	0 Mid
Medium			
1	No ...	Divorced	3 Mid
Medium			
2	No ...	Married	3 Mid
Medium			
3	No ...	Single	2 Mid
Small			
4	Yes ...	Divorced	0 Senior
Medium			

Remote Work	Leadership Opportunities	Innovation Opportunities \
0	No	No
1	No	No
2	No	No
3	Yes	No
4	No	No

Company Reputation	Employee Recognition	Attrition
0	Excellent	Medium Stayed
1	Fair	Low Stayed
2	Poor	Low Stayed
3	Good	Medium Stayed
4	Fair	Medium Stayed

[5 rows x 22 columns]

3. Train-Validation Split [5 marks]

3.1 Import required libraries

```
# Import Train Test Split
from sklearn.model_selection import train_test_split
```

3.2 Define feature and target variables [2 Mark]

```
# Put all the feature variables in X
X = emp_df3.drop(columns=['Attrition'])
# Put the target variable in y
y = emp_df3['Attrition']
```

3.3 Split the data [3 Marks]

```
# Split the data into 70% train data and 30% validation data
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3,
random_state=42)
```

```
# Print the shapes of the resulting datasets
print("Training Features Shape:", X_train.shape)
print("Validation Features Shape:", X_val.shape)
print("Training Labels Shape:", y_train.shape)
print("Validation Labels Shape:", y_val.shape)
```

```
Training Features Shape: (49444, 21)
Validation Features Shape: (21191, 21)
Training Labels Shape: (49444,)
Validation Labels Shape: (21191,)
```

4. EDA on training data [20 marks]

4.1 Perform univariate analysis [6 marks]

Perform univariate analysis on training data for all the numerical columns.

4.1.1 Select numerical columns from training data [1 Mark]

```
# Select numerical columns
numerical_columns = X_train.select_dtypes(include=['int64',
'float64']).columns
```

4.1.2 Plot distribution of numerical columns [5 Marks]


```
# Plot all the numerical columns to understand their distribution



# Import necessary libraries
import seaborn as sns
import matplotlib.pyplot as plt

# Select numerical columns
numerical_columns = X_train.select_dtypes(include=['int64',
'float64']).columns


# Plot all the numerical columns to understand their distribution
for column in numerical_columns:
```

```
plt.figure(figsize=(8, 5))
sns.histplot(X_train[column], kde=True, bins=30, color='blue') #
Histogram with a KDE overlay
plt.title(f"Distribution of {column}")
plt.xlabel(column)
plt.ylabel("Frequency")
plt.show()
```







Distribution of Number of Promotions



Distribution of Distance from Home






4.2 Perform correlation analysis [4 Marks]

Check the correlation among different numerical variables.

```
# Create correlation matrix for numerical columns
correlation_matrix = X_train[numerical_columns].corr()


# Plot Heatmap of the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm",
            fmt=".2f")
plt.title("Correlation Matrix Heatmap")
plt.show()
```



4.3 Check class balance [2 Marks]

Check the distribution of target variable in training set to check class balance.

```
# Plot a bar chart to check class balance
plt.figure(figsize=(8, 5))
sns.countplot(x=y_train, palette="coolwarm")
plt.title("Distribution of Target Variable (Class Balance)")
plt.xlabel("Classes")
plt.ylabel("Frequency")
plt.show()
```

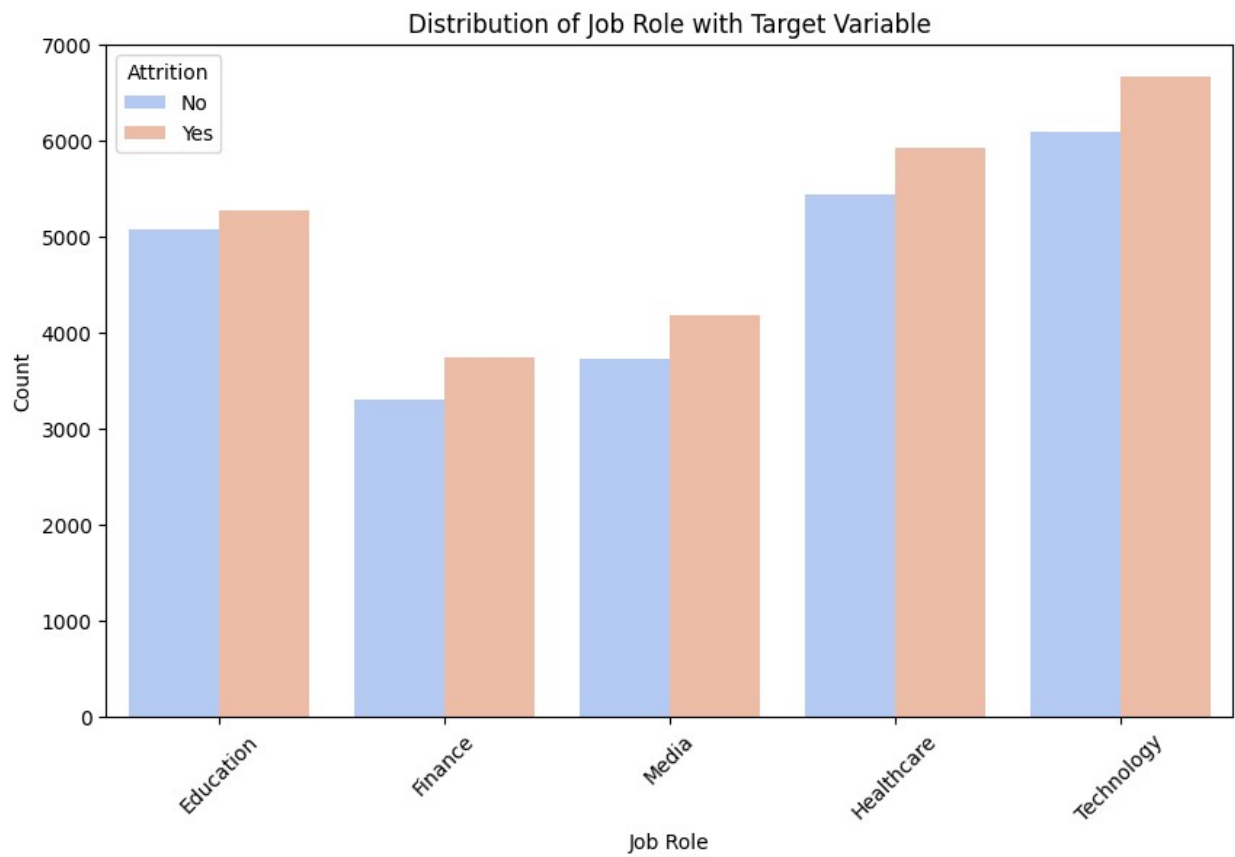
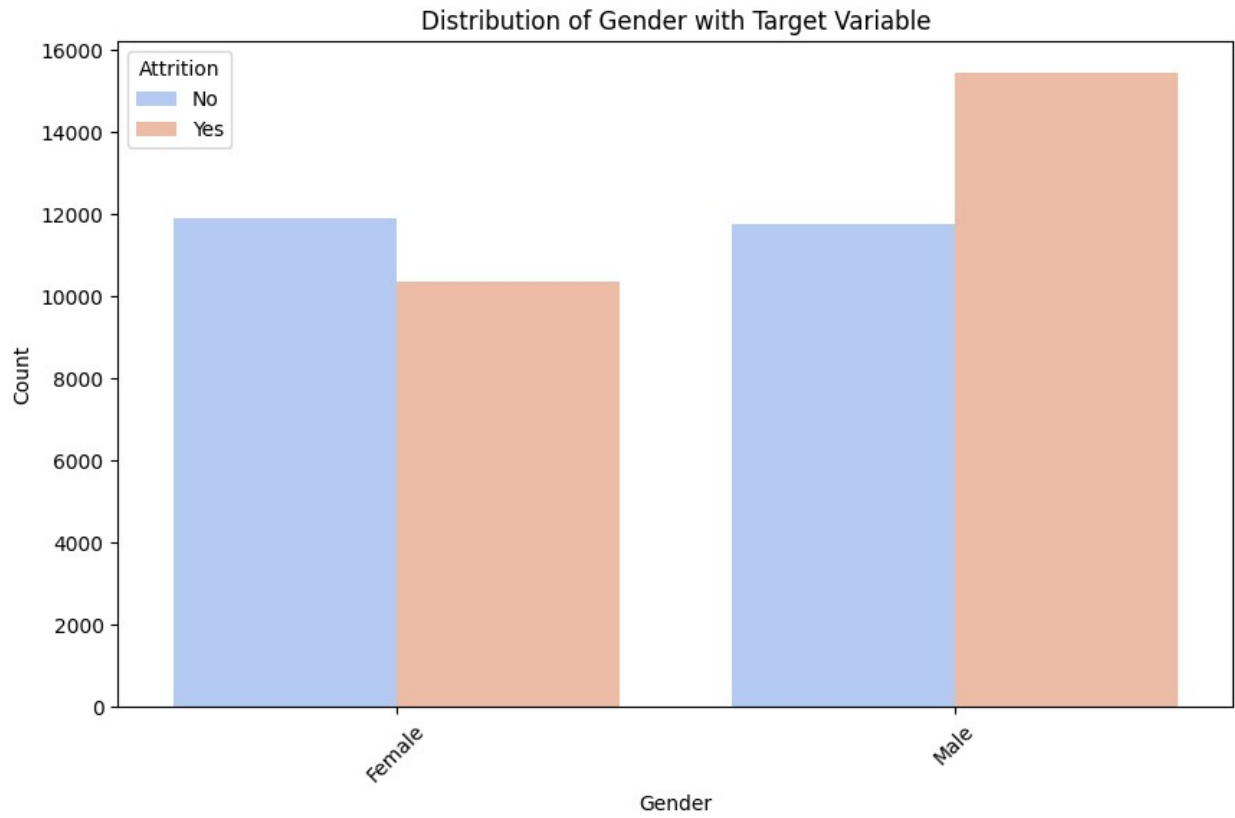



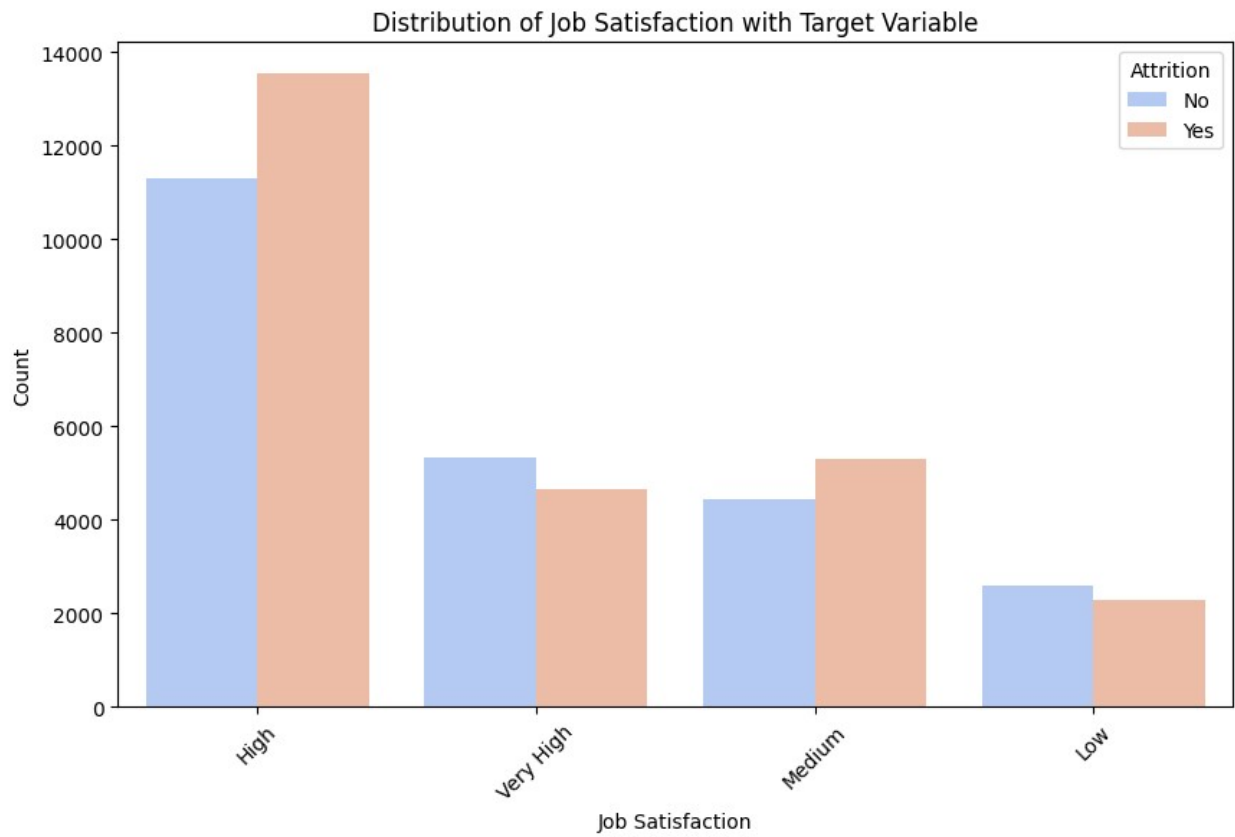

4.4 Perform bivariate analysis [8 Marks]

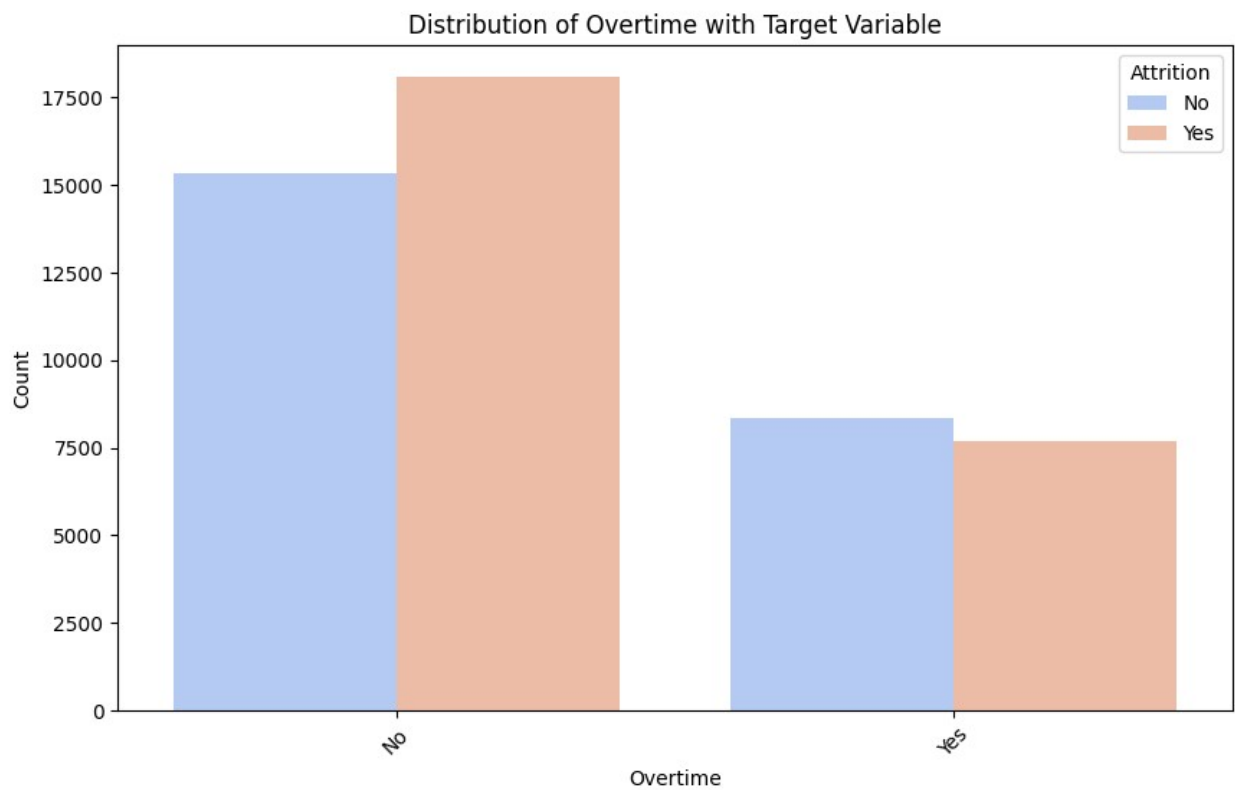
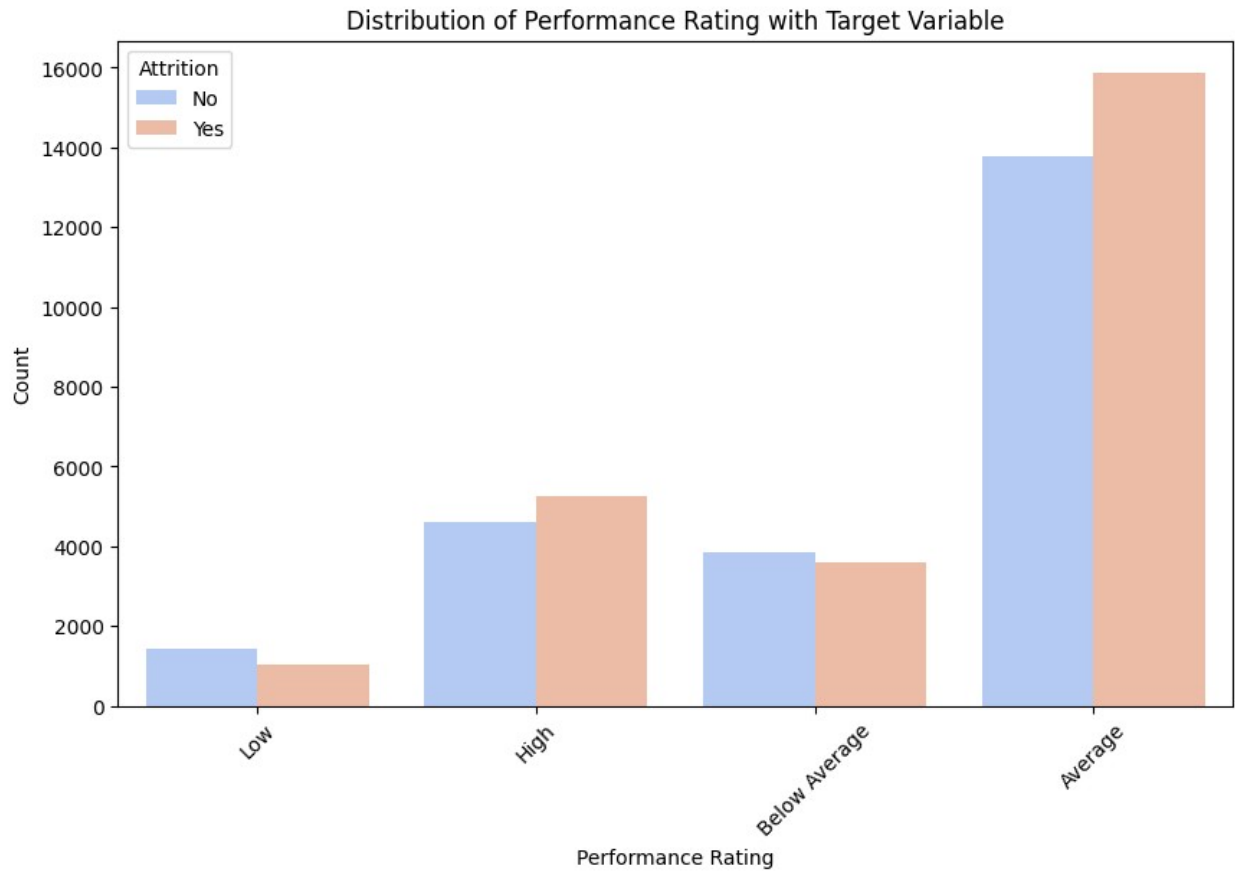
Perform bivariate analysis on training data between all the categorical columns and target variable to analyse how the categorical variables influence the target variable.

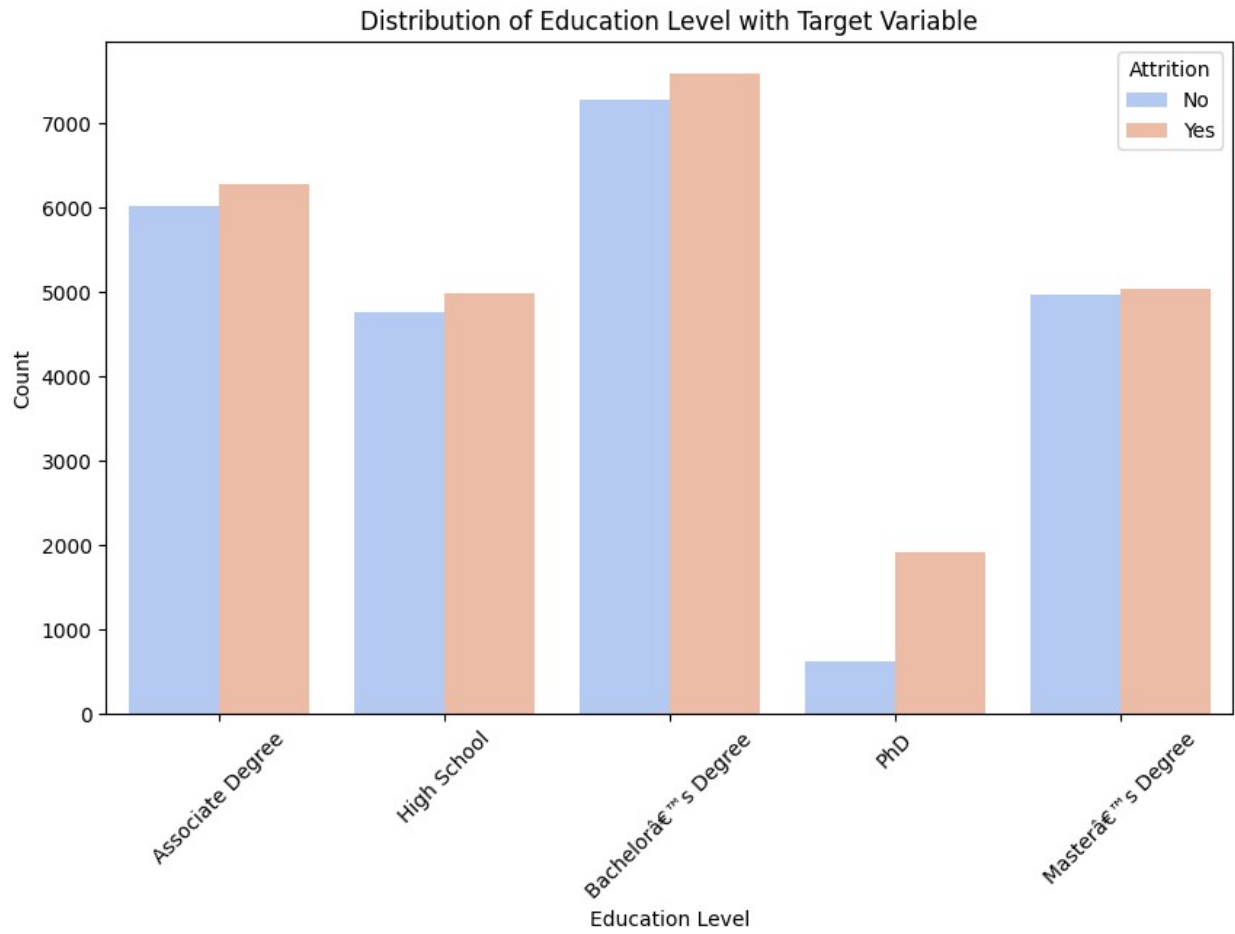
```
# Plot distribution for each categorical column with target variable
# Select categorical columns from the training data
categorical_columns =
X_train.select_dtypes(include=['object']).columns

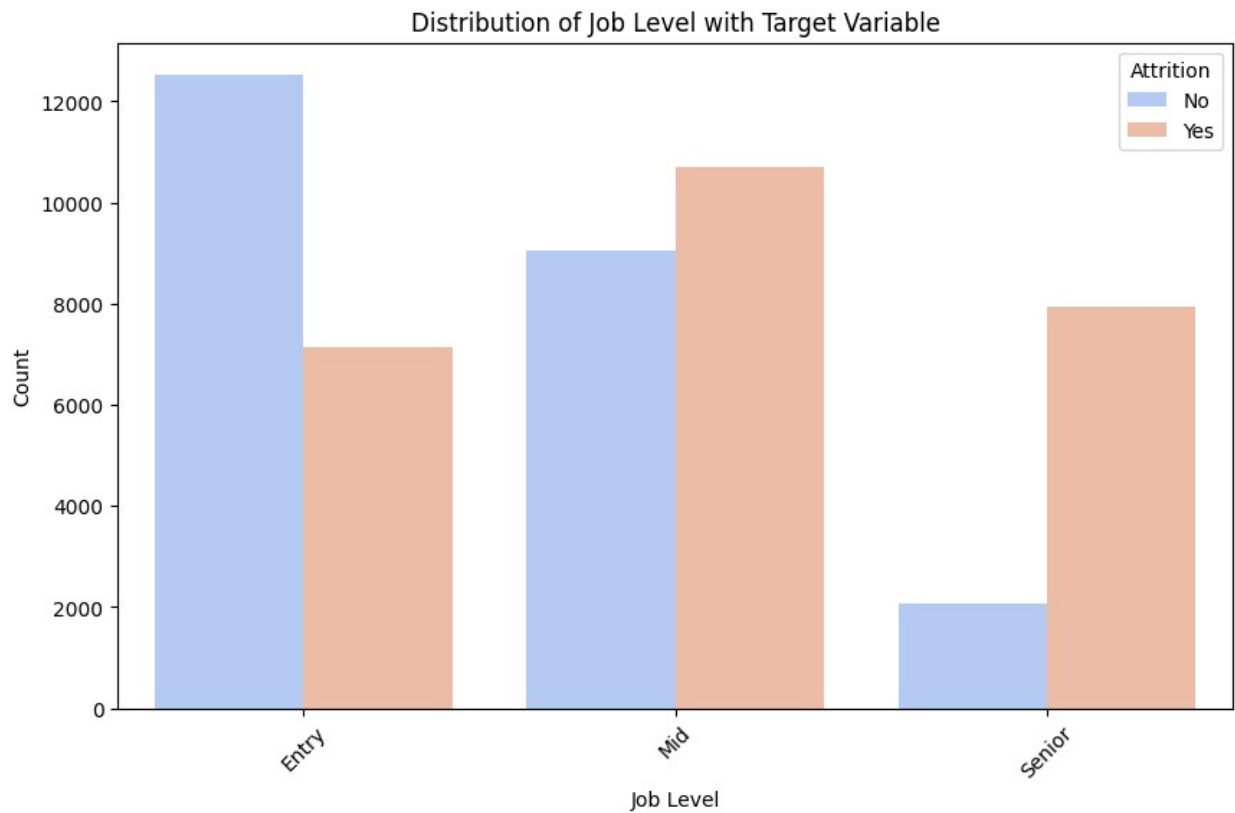
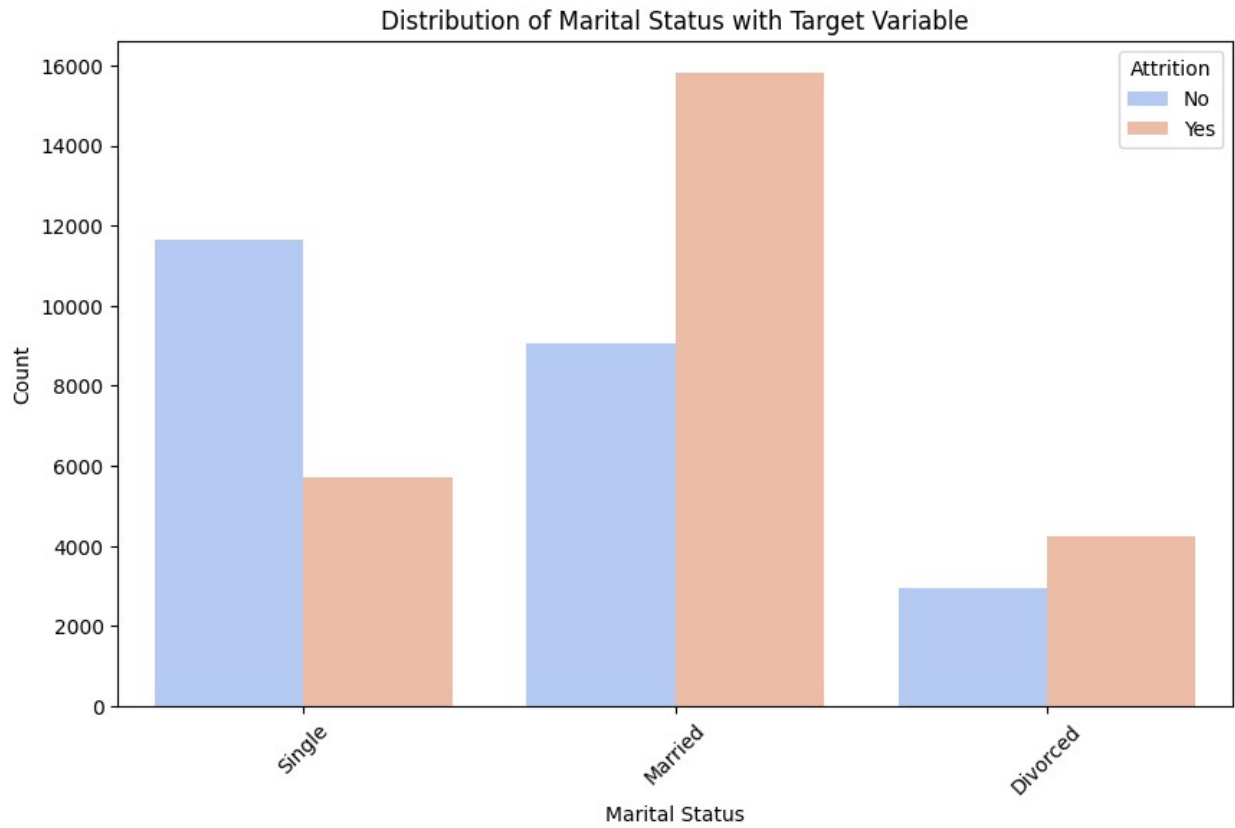
# Plot the distribution for each categorical column with the target
variable
for column in categorical_columns:
    plt.figure(figsize=(10, 6))
    sns.countplot(x=X_train[column], hue=y_train, palette="coolwarm")
    plt.title(f"Distribution of {column} with Target Variable")
    plt.xlabel(column)
    plt.ylabel("Count")
    plt.legend(title="Attrition", labels=["No", "Yes"])
    plt.xticks(rotation=45)
    plt.show()
```

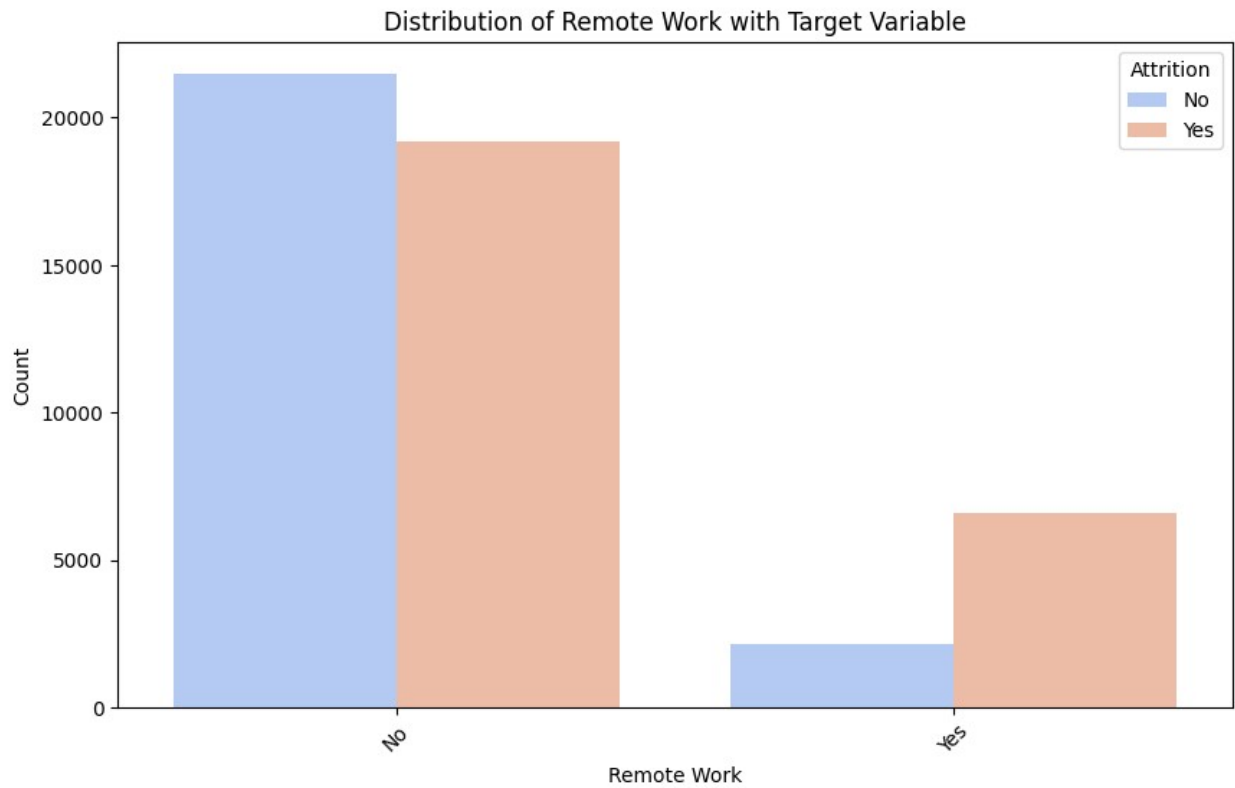
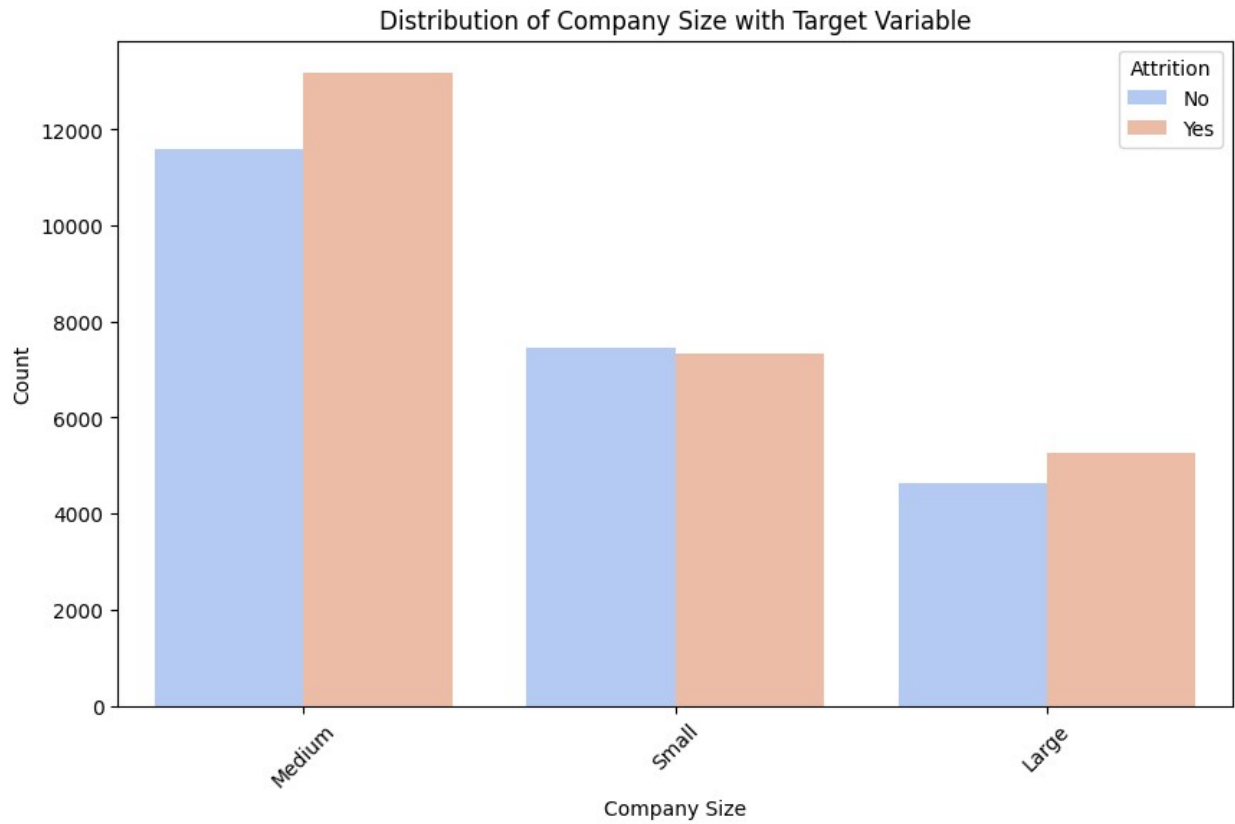



















5. EDA on validation data [OPTIONAL]

5.1 Perform univariate analysis


Perform univariate analysis on validation data for all the numerical columns.



5.1.1 Select numerical columns from validation data


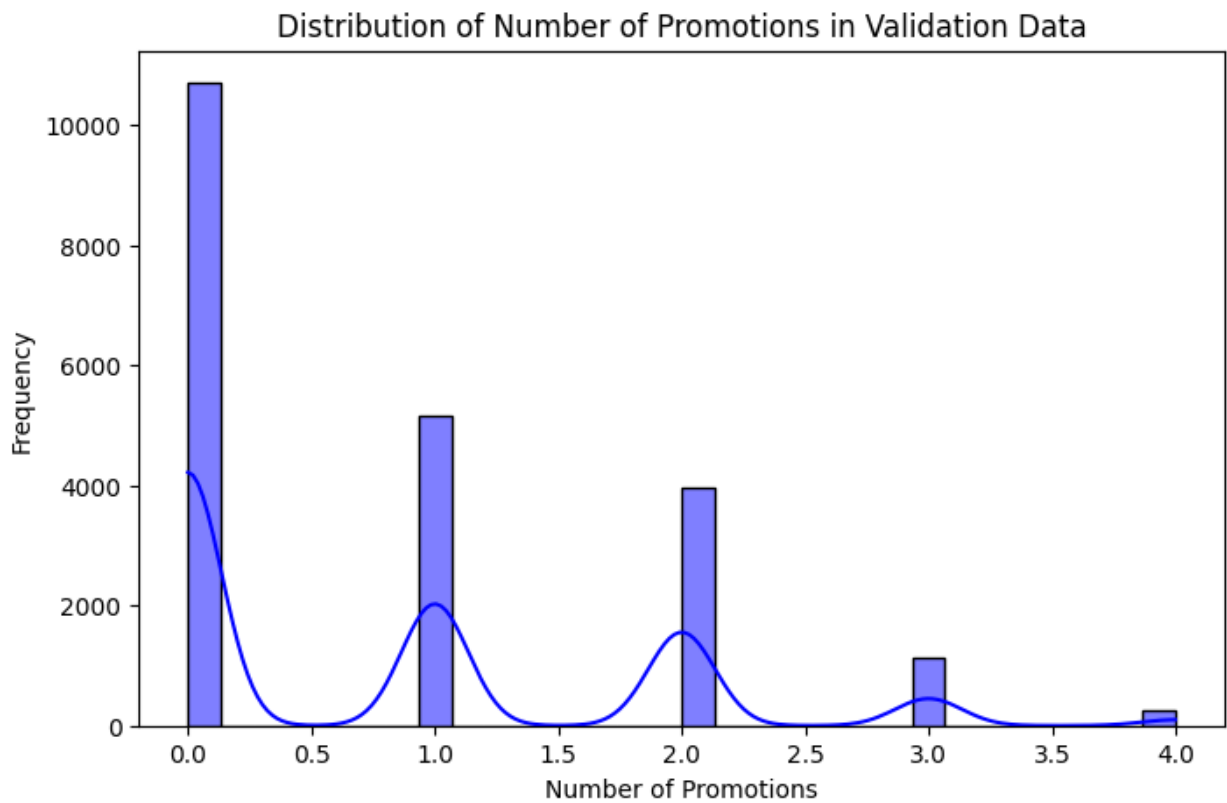
```
# Select numerical columns
numerical_columns = X_val.select_dtypes(include=['int64',
'float64']).columns
```

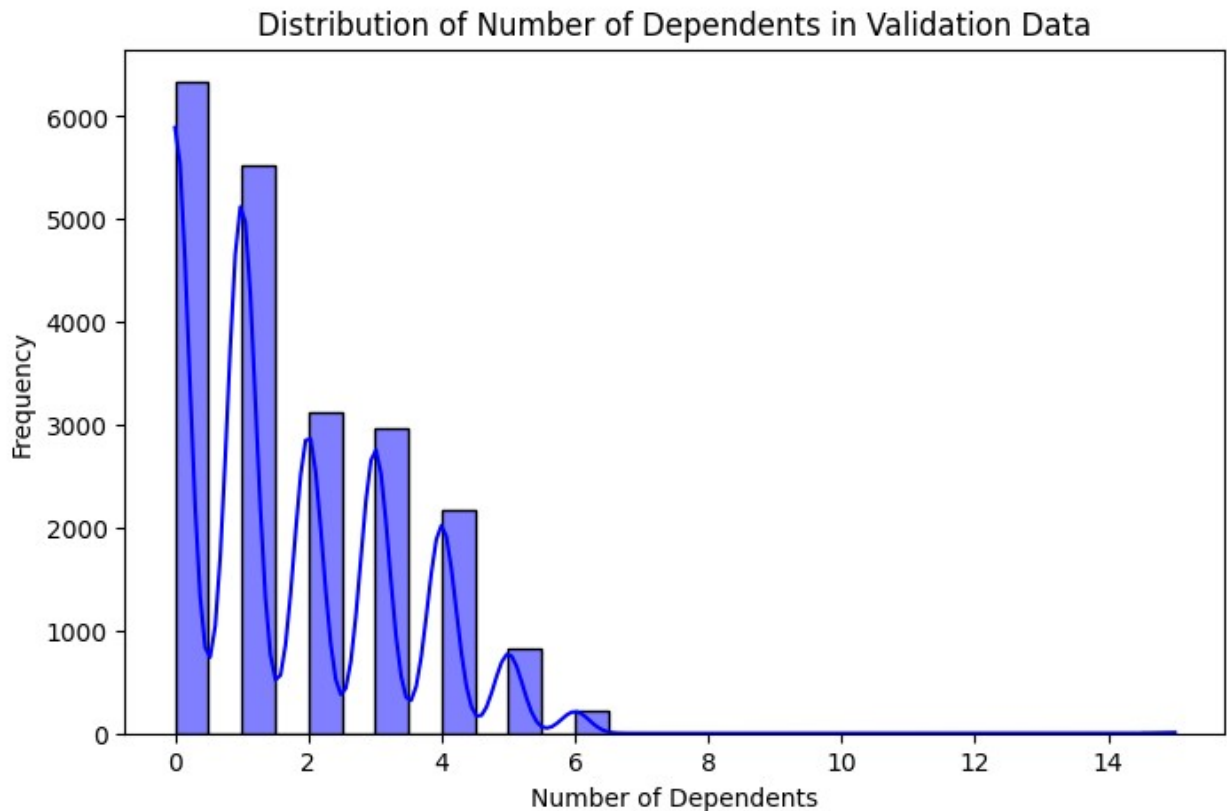
5.1.2 Plot distribution of numerical columns

```
# Plot all the numerical columns to understand their distribution
for column in numerical_columns:
    plt.figure(figsize=(8, 5))
    sns.histplot(X_val[column], kde=True, bins=30, color='blue') #
    Histogram with KDE
    plt.title(f"Distribution of {column} in Validation Data")
    plt.xlabel(column)
    plt.ylabel("Frequency")
    plt.show()
```










5.2 Perform correlation analysis

Check the correlation among different numerical variables.


```
# Create correlation matrix for numerical columns
correlation_matrix = X_val[numerical_columns].corr()
# Plot Heatmap of the correlation matrix
plt.figure(figsize=(10, 8)) # Adjust figure size if required
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm",
            fmt=".2f", linewidths=0.5)
plt.title("Correlation Matrix Heatmap (Validation Data)")
plt.show()
```



5.3 Check class balance

Check the distribution of target variable in validation data to check class balance.

```
# Plot a bar chart to check class balance
plt.figure(figsize=(8, 5))
sns.countplot(x=y_val, palette="coolwarm")
plt.title("Class Balance in Validation Data")
plt.xlabel("Classes")
plt.ylabel("Frequency")
plt.show()
```





5.4 Perform bivariate analysis


Perform bivariate analysis on validation data between all the categorical columns and target variable to analyse how the categorical variables influence the target variable.


```
# Plot distribution for each categorical column with target variable
# Select categorical columns from the validation data
categorical_columns = X_val.select_dtypes(include=['object']).columns


# Plot the distribution of each categorical column with the target
variable
for column in categorical_columns:
    plt.figure(figsize=(10, 6))
    sns.countplot(x=X_val[column], hue=y_val, palette="coolwarm")
    plt.title(f"Distribution of {column} with Target Variable")
    plt.xlabel(column)
    plt.ylabel("Count")
    plt.legend(title="Attrition", labels=["No", "Yes"])
    plt.xticks(rotation=45) # Adjust for better readability
    plt.show()
```


































6. Feature Engineering [20 marks]

6.1 Dummy variable creation [15 marks]

The next step is to deal with the categorical variables present in the data.

6.1.1 Identify categorical columns where dummy variables are required [1 Mark]

```
# Check the categorical columns
categorical_columns =
X_train.select_dtypes(include=['object']).columns
```

```
# Output the categorical columns
print("Categorical columns in the dataset:")
print(categorical_columns)
```

```
Categorical columns in the dataset:
Index(['Gender', 'Job Role', 'Work-Life Balance', 'Job Satisfaction',
      'Performance Rating', 'Overtime', 'Education Level', 'Marital
      Status',
      'Job Level', 'Company Size', 'Remote Work', 'Leadership
      Opportunities',
      'Innovation Opportunities', 'Company Reputation',
```

```
'Employee Recognition'],  
dtype='object')
```

6.1.2 Create dummy variables for independent columns in training set [3 Marks]

```
# Create dummy variables using the 'get_dummies' for independent  
columns  
X_train_encoded = pd.get_dummies(X_train, drop_first=True) # Drop  
first category to prevent multicollinearity  
X_val_encoded = pd.get_dummies(X_val, drop_first=True)  
  
# Add the results to the master DataFrame  
  
# Check the resulting shapes of the master DataFrames  
print("Shape of Training DataFrame after encoding:",  
X_train_encoded.shape)  
print("Shape of Validation DataFrame after encoding:",  
X_val_encoded.shape)
```

```
Shape of Training DataFrame after encoding: (49444, 40)  
Shape of Validation DataFrame after encoding: (21191, 40)
```

Now, drop the original categorical columns and check the DataFrame

```
# Drop the original categorical columns and check the DataFrame  
# Drop original categorical columns from training and validation data  
X_train_encoded = X_train_encoded.drop(columns=categorical_columns)  
X_val_encoded = X_val_encoded.drop(columns=categorical_columns)  
  
# Check the resulting DataFrame  
print("Training DataFrame after dropping categorical columns:")  
print(X_train_encoded.head())  
  
print("\nValidation DataFrame after dropping categorical columns:")  
print(X_val_encoded.head())
```

Training DataFrame after dropping categorical columns:

	Age	Years at Company	Monthly Income	Number of Promotions	\
41465	46	26	4617	2	
69350	43	26	4447	3	
28247	27	8	9762	0	
3217	47	29	5244	2	
73636	25	10	4455	0	

	Distance from Home	Number of Dependents
41465	59.0	2
69350	64.0	0
28247	84.0	1
3217	59.0	1
73636	36.0	0

Validation DataFrame after dropping categorical columns:

	Age	Years at Company	Monthly Income	Number of Promotions	\
23813	46	1	7740	3	
14537	32	6	8779	1	
45192	53	10	11683	1	
13765	50	2	7305	2	
3411	48	24	10114	0	

	Distance from Home	Number of Dependents
23813	78.0	2
14537	91.0	0
45192	71.0	1
13765	71.0	1
3411	2.0	1

6.1.3 Create dummy variables for independent columns in validation set [3 Marks]

```
# Create dummy variables using the 'get_dummies' for independent
columns

# Create dummy variables for categorical columns in training and
validation sets
X_train_encoded = pd.get_dummies(X_train, drop_first=True) # Training
data
X_val_encoded = pd.get_dummies(X_val, drop_first=True)      #
Validation data

# Display the resulting shapes after encoding
print("Shape of Training DataFrame after encoding:",
X_train_encoded.shape)
print("Shape of Validation DataFrame after encoding:",
X_val_encoded.shape)

# Add the results to the master DataFrame

Shape of Training DataFrame after encoding: (49444, 40)
Shape of Validation DataFrame after encoding: (21191, 40)
```

Now, drop the original categorical columns and check the DataFrame

```
# Drop categorical columns and check the DataFrame
X_train_encoded = X_train.drop(columns=categorical_columns)
X_val_encoded = X_val.drop(columns=categorical_columns)

# Check the resulting DataFrame
print("Training DataFrame after dropping categorical columns:")
print(X_train_encoded.head())
```

```
print("\nValidation DataFrame after dropping categorical columns:")
print(X_val_encoded.head())
```

Training DataFrame after dropping categorical columns:

	Age	Years at Company	Monthly Income	Number of Promotions	\
41465	46	26	4617	2	
69350	43	26	4447	3	
28247	27	8	9762	0	
3217	47	29	5244	2	
73636	25	10	4455	0	

	Distance from Home	Number of Dependents
41465	59.0	2
69350	64.0	0
28247	84.0	1
3217	59.0	1
73636	36.0	0

Validation DataFrame after dropping categorical columns:

	Age	Years at Company	Monthly Income	Number of Promotions	\
23813	46	1	7740	3	
14537	32	6	8779	1	
45192	53	10	11683	1	
13765	50	2	7305	2	
3411	48	24	10114	0	

	Distance from Home	Number of Dependents
23813	78.0	2
14537	91.0	0
45192	71.0	1
13765	71.0	1
3411	2.0	1

6.1.4 Create DataFrame for dependent column in both training and validation set [1 Mark]

```
# Convert y_train and y_validation to DataFrame to create dummy
variables
# Convert y_train and y_val to DataFrames
y_train_df = pd.DataFrame(y_train, columns=['Attrition'])
y_val_df = pd.DataFrame(y_val, columns=['Attrition'])

# Create dummy variables for the target variable
y_train_encoded = pd.get_dummies(y_train_df, drop_first=True)
y_val_encoded = pd.get_dummies(y_val_df, drop_first=True)

# Check the encoded target variable
print("Encoded y_train:")
print(y_train_encoded.head())
```

```
print("\nEncoded y_val:")
print(y_val_encoded.head())
```

```
Encoded y_train:
      Attrition Stayed
41465          False
69350          False
28247          False
3217           False
73636          False
```

```
Encoded y_val:
      Attrition Stayed
23813          True
14537          False
45192          False
13765          False
3411           True
```

6.1.5 Create dummy variables for dependent column in training set [3 Marks]

```
# Create dummy variables using the 'get_dummies' for dependent column
y_train_encoded = pd.get_dummies(y_train, drop_first=True)
y_val_encoded = pd.get_dummies(y_val, drop_first=True)
```

```
# Check the resulting dummy variables
```

```
print("Encoded y_train:")
print(y_train_encoded.head())
```

```
print("\nEncoded y_val:")
print(y_val_encoded.head())
```

```
Encoded y_train:
      Stayed
41465  False
69350  False
28247  False
3217   False
73636  False
```

```
Encoded y_val:
      Stayed
23813  True
14537  False
45192  False
13765  False
3411   True
```

6.1.6 Create dummy variable for dependent column in validation set [3 Marks]

```
# Create dummy variables using the 'get_dummies' for dependent column
y_val_encoded = pd.get_dummies(y_val, drop_first=True)
```

6.1.7 Drop redundant columns [1 Mark]

```
# Drop redundant columns from both train and validation
def drop_redundant_columns(X_train, X_val):
    # Identify constant columns
    constant_columns = [col for col in X_train.columns if
X_train[col].nunique() == 1]

    # Identify duplicate columns
    duplicate_columns =
X_train.T[X_train.T.duplicated()].index.tolist()

    # Identify columns in train but not in validation
    train_only_columns = [col for col in X_train.columns if col not in
X_val.columns]

    # Consolidate all columns to drop
    columns_to_drop = list(set(constant_columns + duplicate_columns +
train_only_columns))

    # Drop from both datasets
    X_train.drop(columns=columns_to_drop, inplace=True)
    X_val.drop(columns=columns_to_drop, inplace=True)

    return X_train, X_val

# Apply to datasets
X_train, X_val = drop_redundant_columns(X_train, X_val)
```

6.2 Feature scaling [5 marks]

Apply feature scaling to the numeric columns to bring them to a common range and ensure consistent scaling.

6.2.1 Import required libraries [1 Mark]

```
# Import the necessary scaling tool from scikit-learn
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

6.2.2 Scale the numerical features [4 Marks]

```
numeric_cols = X_train.select_dtypes(include=['int64',
'float64']).columns

# Initialize StandardScaler
scaler = StandardScaler()
```



```
# Scale the numeric features present in the training set
X_train_scaled =
pd.DataFrame(scaler.fit_transform(X_train[numeric_cols]),
columns=numeric_cols)

# Scale the numerical features present in the validation set
X_val_scaled = pd.DataFrame(scaler.transform(X_val[numeric_cols]),
columns=numeric_cols)
```

7. Model Building [40 marks]

7.1 Feature selection [5 marks]

As there are a lot of variables present in the data, Recursive Feature Elimination (RFE) will be used to select the most influential features for building the model.

7.1.1 Import required libraries [1 Mark]

```
# Import 'LogisticRegression' and create a LogisticRegression object
from sklearn.linear_model import LogisticRegression
# Create a Logistic Regression object
log_reg = LogisticRegression()
```

7.1.2 Import RFE and select 15 variables [3 Mark]

```
# Import RFE and select 15 variables
from sklearn.feature_selection import RFE
from sklearn.datasets import make_classification
# Generate synthetic dataset
X, y = make_classification(n_samples=200, n_features=20,
n_informative=15, n_redundant=5, random_state=42)
X = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 21)])
y = pd.Series(y, name='Target')

# Display the features selected by RFE
# Initialize Logistic Regression model
log_reg = LogisticRegression(random_state=42)

# Perform RFE
rfe = RFE(estimator=log_reg, n_features_to_select=15)
rfe.fit(X, y)

RFE(estimator=LogisticRegression(random_state=42),
n_features_to_select=15)
```

7.1.3 Store the selected features [1 Mark]

```
# Put columns selected by RFE into variable 'col'
col = X.columns[rfe.support_]
```

7.2 Building Logistic Regression Model [20 marks]

Now that you have selected the variables through RFE, use these features to build a logistic regression model with statsmodels. This will allow you to assess the statistical aspects, such as p-values and VIFs, which are important for checking multicollinearity and ensuring that the predictors are not highly correlated with each other, as this could distort the model's coefficients.

7.2.1 Select relevant columns on training set [1 Mark]

```
# Select only the columns selected by RFE
selected_columns = X.columns[rfe.support_]
X_selected = X[selected_columns]

# View the training data
X, y = make_classification(n_samples=200, n_features=20,
random_state=42)
X = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(1, 21)])
y = pd.Series(y, name="Target")
# Step 2: Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

7.2.2 Add constant to training set [1 Mark]

```
# Import statsmodels and add constant to training set
import statsmodels.api as sm
X_train_sm = sm.add_constant(X_train)
```

7.2.3 Fit logistic regression model [3 Marks]

```
# Fit a logistic regression model on X_train after adding a constant
and output the summary
logit_model = sm.Logit(y_train, X_train_sm) # Initialize the model
result = logit_model.fit() # Fit the model

Optimization terminated successfully.
    Current function value: 0.328455
    Iterations 23
```

Model Interpretation

The output summary table will provide the features used for building model along with coefficient of each of the feature and their p-value. The p-value in a logistic regression model is used to assess the statistical significance of each coefficient. Lesser the p-value, more significant the feature is in the model.

A positive coefficient will indicate that an increase in the value of feature would increase the odds of the event occurring. On the other hand, a negative coefficient means the opposite, i.e, an increase in the value of feature would decrease the odds of the event occurring.

7.2.4 Evaluate VIF of features [3 Marks]

```
# Import 'variance_inflation_factor'
from statsmodels.stats.outliers_influence import
variance_inflation_factor

# Make a VIF DataFrame for all the variables present
vif_data = pd.DataFrame()
vif_data["Feature"] = X.columns # All feature names
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]
```

Proceed to the next step if p-values and VIFs are within acceptable ranges. If you observe high p-values or VIFs, create new cells to drop the features and retrain the model.

7.2.5 Make predictions on training set [2 Marks]

```
# Predict the probabilities on the training set
y_train_prob = result.predict(X_train_sm)
```

7.2.6 Format the prediction output [1 Mark]

```
# Reshape it into an array
y_train_prob_array = y_train_prob.to_numpy()
y_train_prob_resaped = y_train_prob_array.reshape(-1, 1)
```

7.2.7 Create a DataFrame with the actual stayed flag and the predicted probabilities [1 Mark]

```
# Create a new DataFrame containing the actual stayed flag and the
probabilities predicted by the model
results_df = pd.DataFrame({
    "Actual_Stayed": y_train, # Actual target values from
training set
    "Predicted_Probability": y_train_prob # Predicted probabilities
from the model
})
```

7.2.8 Create a new column 'Predicted' with 1 if predicted probabilities are greater than 0.5 else 0 [1 Mark]

```
# Create a new column 'Predicted' with 1 if predicted probabilities
are greater than 0.5 else 0
results_df["Predicted"] = (results_df["Predicted_Probability"] >
0.5).astype(int)
```

Evaluation of performance of Model

Evaluate the performance of the model based on the predictions made on the training set.

7.2.9 Check the accuracy of the model based on the predictions made on the training set [1 Mark]

```
# Import metrics from sklearn for evaluation
from sklearn import metrics

# Check the overall accuracy
accuracy = metrics.accuracy_score(results_df["Actual_Stayed"],
results_df["Predicted"])
```

7.2.10 Create a confusion matrix based on the predictions made on the training set [1 mark]

```
# Create confusion matrix
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(results_df["Actual_Stayed"],
results_df["Predicted"])
```

7.2.11 Create variables for true positive, true negative, false positive and false negative [1 Mark]

```
# Create variables for true positive, true negative, false positive
and false negative
# Extract values
TN = conf_matrix[0, 0] # True Negatives
FP = conf_matrix[0, 1] # False Positives
FN = conf_matrix[1, 0] # False Negatives
TP = conf_matrix[1, 1] # True Positives

accuracy = (TP + TN) / (TP + TN + FP + FN)
print(f"\nAccuracy: {accuracy:.2f}")
```

Accuracy: 0.88

7.2.12 Calculate sensitivity and specificity of model [2 Marks]

```
# Calculate sensitivity
sensitivity = TP / (TP + FN)

# Display sensitivity
print(f"\nSensitivity (Recall): {sensitivity:.2f}")
```

Sensitivity (Recall): 0.90

```
# Calculate specificity
specificity = TN / (TN + FP)
print(f"\nSpecificity (True Negative Rate): {specificity:.2f}")
```

Specificity (True Negative Rate): 0.87

7.2.13 Calculate precision and recall of model [2 Marks]

```
# Calculate precision
```

```
precision = TP / (TP + FP)
```

```
print(f"\nPrecision (Positive Predictive Value): {precision:.2f}")
```

Precision (Positive Predictive Value): 0.86

```
# Calculate recall
```

```
recall = TP / (TP + FN)
```

```
print(f"\nRecall (Sensitivity/True Positive Rate): {recall:.2f}")
```

Recall (Sensitivity/True Positive Rate): 0.90

7.3 Find the optimal cutoff [15 marks]

Find the optimal cutoff to improve model performance. While a default threshold of 0.5 was used for initial evaluation, optimising this threshold can enhance the model's performance.

First, plot the ROC curve and check AUC.

7.3.1 Plot ROC curve [3 Marks]

```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
```

```
# Define ROC function
```

```
def plot_roc_curve(y_true, y_prob):
```

```
    # Compute False Positive Rate (FPR), True Positive Rate (TPR), and thresholds
```

```
    fpr, tpr, thresholds = roc_curve(y_true, y_prob)
```

```
    # Compute AUC
```

```
    auc = roc_auc_score(y_true, y_prob)
```

```
    # Plot ROC curve
```

```
    plt.figure(figsize=(10, 6))
```

```
    plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc:.2f})",
             color="blue")
```

```
    plt.plot([0, 1], [0, 1], linestyle="--", color="grey",
             label="Random Classifier")
```


```

plt.xlabel("False Positive Rate (FPR)")
plt.ylabel("True Positive Rate (TPR)")
plt.title("Receiver Operating Characteristic (ROC) Curve")
plt.legend(loc="lower right")
plt.grid()
plt.show()

print(f"AUC Score: {auc:.2f}")
return fpr, tpr, thresholds

# Call the ROC function
fpr, tpr, thresholds = plot_roc_curve(y_train, y_train_prob)

```



AUC Score: 0.94

Sensitivity and Specificity tradeoff

Check sensitivity and specificity tradeoff to find the optimal cutoff point.

7.3.2 Predict on training set at various probability cutoffs [1 Mark]

```

# Predict on training data by creating columns with different
# probability cutoffs to explore the impact of cutoff on model
# performance
from sklearn.metrics import confusion_matrix, accuracy_score,

```

```

precision_score, recall_score

# Function to add columns for cutoff-based predictions
def explore_cutoffs(y_true, y_prob, cutoffs):

    metrics = []

    for cutoff in cutoffs:
        # Classify based on the current cutoff
        y_pred = (y_prob >= cutoff).astype(int)

        # Calculate metrics
        accuracy = accuracy_score(y_true, y_pred)
        precision = precision_score(y_true, y_pred, zero_division=0)
# Avoid division by zero
        recall = recall_score(y_true, y_pred, zero_division=0)

        # Store the metrics for this cutoff
        metrics.append({"Cutoff": cutoff, "Accuracy": accuracy,
"Precision": precision, "Recall": recall})

    # Create a DataFrame to summarize the results
    metrics_df = pd.DataFrame(metrics)

    return metrics_df

```

7.3.3 Plot for accuracy, sensitivity, specificity at different probability cutoffs [2 Marks]

```

# Create a DataFrame to see the values of accuracy, sensitivity, and
specificity at different values of probability cutoffs

import pandas as pd
from sklearn.metrics import confusion_matrix, accuracy_score

def evaluate_metrics_cutoffs(y_true, y_prob, cutoffs):

    metrics = []

    for cutoff in cutoffs:
        # Classify predictions based on the current cutoff
        y_pred = (y_prob >= cutoff).astype(int)

        # Compute confusion matrix
        conf_matrix = confusion_matrix(y_true, y_pred)
        TN = conf_matrix[0, 0] # True Negatives
        FP = conf_matrix[0, 1] # False Positives
        FN = conf_matrix[1, 0] # False Negatives
        TP = conf_matrix[1, 1] # True Positives

        # Calculate metrics

```

```

        accuracy = accuracy_score(y_true, y_pred)
        sensitivity = TP / (TP + FN) if (TP + FN) > 0 else 0 # Recall
        (sensitivity)
        specificity = TN / (TN + FP) if (TN + FP) > 0 else 0 #
        Specificity

        # Append results
        metrics.append({"Cutoff": cutoff, "Accuracy": accuracy,
        "Sensitivity": sensitivity, "Specificity": specificity})

        # Create DataFrame
        metrics_df = pd.DataFrame(metrics)

        return metrics_df

```

Example Workflow

Step 1: Define cutoffs to explore

```
cutoffs = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
```

Step 2: Use function and generate metrics DataFrame

```
metrics_df = evaluate_metrics_cutoffs(y_train, y_train_prob, cutoffs)
```

Step 3: Display metrics DataFrame

```
print("\nMetrics at Different Probability Cutoffs:")
```

```
print(metrics_df)
```

Metrics at Different Probability Cutoffs:

	Cutoff	Accuracy	Sensitivity	Specificity
0	0.1	0.80000	0.987179	0.621951
1	0.2	0.81875	0.974359	0.670732
2	0.3	0.85000	0.961538	0.743902
3	0.4	0.86875	0.935897	0.804878
4	0.5	0.88125	0.897436	0.865854
5	0.6	0.85000	0.807692	0.890244
6	0.7	0.82500	0.717949	0.926829
7	0.8	0.78125	0.602564	0.951220
8	0.9	0.71250	0.448718	0.963415

Plot accuracy, sensitivity, and specificity at different values of probability cutoffs

```
from sklearn.metrics import confusion_matrix, accuracy_score
```

```
def evaluate_metrics_cutoffs(y_true, y_prob, cutoffs):
```

```
    metrics = []
```

```
    for cutoff in cutoffs:
```

```
        # Classify predictions based on the current cutoff
```

```
        y_pred = (y_prob >= cutoff).astype(int)
```



```

# Compute confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred)
TN = conf_matrix[0, 0] # True Negatives
FP = conf_matrix[0, 1] # False Positives
FN = conf_matrix[1, 0] # False Negatives
TP = conf_matrix[1, 1] # True Positives

# Calculate metrics
accuracy = accuracy_score(y_true, y_pred)
sensitivity = TP / (TP + FN) if (TP + FN) > 0 else 0 # Recall
(sensitivity)
specificity = TN / (TN + FP) if (TN + FP) > 0 else 0 #
Specificity

# Append results
metrics.append({"Cutoff": cutoff, "Accuracy": accuracy,
"Sensitivity": sensitivity, "Specificity": specificity})

# Create DataFrame
metrics_df = pd.DataFrame(metrics)

return metrics_df

# Example Workflow
# Step 1: Define cutoffs to explore
cutoffs = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]

# Step 2: Use function and generate metrics DataFrame
metrics_df = evaluate_metrics_cutoffs(y_train, y_train_prob, cutoffs)

# Step 3: Display metrics DataFrame
print("\nMetrics at Different Probability Cutoffs:")
print(metrics_df)

```

```

Metrics at Different Probability Cutoffs:
   Cutoff  Accuracy  Sensitivity  Specificity
0    0.1    0.80000    0.987179    0.621951
1    0.2    0.81875    0.974359    0.670732
2    0.3    0.85000    0.961538    0.743902
3    0.4    0.86875    0.935897    0.804878
4    0.5    0.88125    0.897436    0.865854
5    0.6    0.85000    0.807692    0.890244
6    0.7    0.82500    0.717949    0.926829
7    0.8    0.78125    0.602564    0.951220
8    0.9    0.71250    0.448718    0.963415

```

7.3.4 Create a column for final prediction based on the optimal cutoff [2 Marks]

```
# Create a column for final prediction based on the optimal cutoff
y_train_prob = result.predict(X_train_sm)
results_df = pd.DataFrame({"Actual": y_train, "Predicted_Probability":
y_train_prob})
optimal_cutoff = 0.4
results_df["Final_Prediction"] = (results_df["Predicted_Probability"]
> optimal_cutoff).astype(int)
```

7.3.5 Calculate model's accuracy [1Mark]

```
# Calculate the accuracy
from sklearn.metrics import confusion_matrix, accuracy_score

conf_matrix = confusion_matrix(results_df["Actual"],
results_df["Final_Prediction"])
TN, FP, FN, TP = conf_matrix.ravel()
accuracy = (TP + TN) / (TP + TN + FP + FN)

# Display results
print("\nConfusion Matrix:")
print(conf_matrix)
print(f"\nAccuracy: {accuracy:.2f}")
```

Confusion Matrix:
[[66 16]
[5 73]]

Accuracy: 0.87

7.3.6 Create confusion matrix [1Mark]

```
# Create the confusion matrix once again
conf_matrix = confusion_matrix(results_df["Actual"],
results_df["Final_Prediction"])
```

7.3.7 Create variables for true positive, true negative, false positive and false negative [1Mark]

```
# Create variables for true positive, true negative, false positive
and false negative
TN, FP, FN, TP = conf_matrix.ravel()
```

7.3.8 Calculate sensitivity and specificity of the model [1Mark]

```
# Calculate Sensitivity
sensitivity = TP / (TP + FN)

# Display sensitivity
print(f"\nSensitivity (Recall/True Positive Rate): {sensitivity:.2f}")
```

Sensitivity (Recall/True Positive Rate): 0.94

```
# Calculate Specificity
```

```
specificity = TN / (TN + FP)
```

```
# Display specificity
```

```
print(f"\nSpecificity (True Negative Rate): {specificity:.2f}")
```

Specificity (True Negative Rate): 0.80

7.3.9 Calculate precision and recall of the model [1Mark]

```
# Calculate Precision
```

```
precision = TP / (TP + FP) if (TP + FP) > 0 else 0
```

```
# Display precision
```

```
print(f"\nPrecision (Positive Predictive Value): {precision:.2f}")
```

Precision (Positive Predictive Value): 0.82

```
# Calculate Recall
```

```
recall = TP / (TP + FN) if (TP + FN) > 0 else 0
```

```
print(f"\nRecall (Sensitivity/True Positive Rate): {recall:.2f}")
```

Recall (Sensitivity/True Positive Rate): 0.94

Precision and Recall tradeoff

Check optimal cutoff value by plotting precision-recall curve, and adjust the cutoff based on the precision and recall tradeoff if required.

```
# Import precision-recall curve function
```

```
from sklearn.metrics import precision_recall_curve
```

```
# Check actual and predicted values from initial model
```

```
print("Actual Class Distribution:\n",  
      results_df["Actual"].value_counts())
```

```
print("\nPredicted Class Distribution:\n",  
      results_df["Final_Prediction"].value_counts())
```

Actual Class Distribution:

Actual

0 82

1 78

Name: count, dtype: int64

Predicted Class Distribution:

Final_Prediction

1 89

0 71

Name: count, dtype: int64

7.3.10 Plot precision-recall curve [2 Marks]

```
# Plot precision-recall curve
from sklearn.metrics import precision_recall_curve, auc
import matplotlib.pyplot as plt

def plot_precision_recall_curve(y_true, y_prob):


    # Compute precision, recall, and thresholds
    precision, recall, thresholds = precision_recall_curve(y_true,
y_prob)

    # Compute AUC for Precision-Recall Curve
    pr_auc = auc(recall, precision)
    print(f"\nPrecision-Recall AUC: {pr_auc:.2f}")

    # Plot the Precision-Recall Curve
    plt.figure(figsize=(10, 6))
    plt.plot(recall, precision, label=f"PR Curve (AUC =
{pr_auc:.2f})", color="blue")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.title("Precision-Recall (PR) Curve")
    plt.legend(loc="lower left")
    plt.grid()
    plt.show()

plot_precision_recall_curve(y_train, y_train_prob)
```

Precision-Recall AUC: 0.91



8. Prediction and Model Evaluation [30 marks]

Use the model from the previous step to make predictions on the validation set with the optimal cutoff. Then evaluate the model's performance using metrics such as accuracy, sensitivity, specificity, precision, and recall.

8.1 Make predictions over validation set [15 marks]

8.1.1 Select relevant features for validation set [2 Marks]

```
# Select the relevant features for validation set
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification
# Step 1: Generate synthetic dataset
X, y = make_classification(n_samples=200, n_features=10,
n_informative=8, random_state=42)
X = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(1, 11)])
y = pd.Series(y, name="Target")

# Step 2: Split the dataset into training and validation sets
X_train, X_validation, y_train, y_validation = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Step 3: Train logistic regression model with selected features
selected_features = ["Feature_1", "Feature_3", "Feature_5",
"Feature_6", "Feature_8"] # Features selected during training
X_train_selected = X_train[selected_features]
X_train_selected = sm.add_constant(X_train_selected) # Add intercept
logit_model = sm.Logit(y_train, X_train_selected)
result = logit_model.fit()

Optimization terminated successfully.
    Current function value: 0.569035
    Iterations 6
```

8.1.2 Add constant to X_validation [2 Marks]

```
# Add constant to X_validation
X_validation_selected = X_validation[selected_features]
X_validation_selected = sm.add_constant(X_validation_selected) # Add
intercept to validation set
```

8.1.3 Make predictions over validation set [3 Marks]

```
# Make predictions on the validation set and store it in the variable
'y_validation_pred'

# Step 5: Make predictions on the validation set
y_validation_prob = result.predict(X_validation_selected) # Predict
probabilities
optimal_cutoff = 0.4 # Use the optimal cutoff determined earlier
y_validation_pred = (y_validation_prob >= optimal_cutoff).astype(int)
# Classify using this cutoff

# Step 6: Evaluate performance using confusion matrix
conf_matrix = confusion_matrix(y_validation, y_validation_pred)
TN, FP, FN, TP = conf_matrix.ravel()

# Step 7: Calculate metrics
accuracy = accuracy_score(y_validation, y_validation_pred)
precision = precision_score(y_validation, y_validation_pred,
zero_division=0)
recall = recall_score(y_validation, y_validation_pred,
zero_division=0) # Sensitivity
specificity = TN / (TN + FP) if (TN + FP) > 0 else 0 # True Negative
Rate (Specificity)

# View predictions
# Step 8: Print the metrics and confusion matrix
print("\nConfusion Matrix:")
print(conf_matrix)
```

```

print(f"\nAccuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall (Sensitivity): {recall:.2f}")
print(f"Specificity: {specificity:.2f}")

```

Confusion Matrix:

```

[[ 7 21]
 [ 2 10]]

```

Accuracy: 0.42
 Precision: 0.32
 Recall (Sensitivity): 0.83
 Specificity: 0.25

8.1.4 Create DataFrame with actual values and predicted values for validation set [5 Marks]

```

# Convert 'y_validation_pred' to a DataFrame 'predicted_probability'
predicted_probability = pd.DataFrame(y_validation_prob,
columns=["Predicted_Probability"])

# Convert 'y_validation' to DataFrame 'actual'
actual = pd.DataFrame(y_validation, columns=["Actual"])

# Remove index from both DataFrames 'actual' and
'predicted_probability' to append them side by side
predicted_probability.reset_index(drop=True, inplace=True)
actual.reset_index(drop=True, inplace=True)
results_df = pd.concat([actual, predicted_probability], axis=1)

```

8.1.5 Predict final prediction based on the cutoff value [3 Marks]

```

# Make predictions on the validation set using the optimal cutoff and
store it in a column 'final_prediction'
optimal_cutoff = 0.4
results_df["final_prediction"] = (results_df["Predicted_Probability"]
>= optimal_cutoff).astype(int)

# Check the DataFrame
results_df.head()

```

	Actual	Predicted_Probability	final_prediction
0	NaN	0.460501	1
1	NaN	0.440745	1
2	NaN	0.037016	0
3	NaN	0.677114	1
4	NaN	0.627497	1

8.2 Calculate accuracy of the model [2 marks]

```
# Calculate the overall accuracy
# Make predictions on the validation set using the optimal cutoff
results_df["final_prediction"] = (results_df["Predicted_Probability"]
>= optimal_cutoff).astype(int)

# Check the updated DataFrame
print("\nUpdated DataFrame with Final Predictions:")
print(results_df.head())
```

Updated DataFrame with Final Predictions:

	Actual	Predicted_Probability	final_prediction
0	NaN	0.460501	1
1	NaN	0.440745	1
2	NaN	0.037016	0
3	NaN	0.677114	1
4	NaN	0.627497	1

8.3 Create confusion matrix and create variables for true positive, true negative, false positive and false negative [5 marks]

```
# Create confusion matrix
results_df = pd.DataFrame({
    "Actual": y_validation,
    "final_prediction": y_validation_pred
})

conf_matrix = confusion_matrix(results_df["Actual"],
results_df["final_prediction"])

# Create variables for true positive, true negative, false positive
and false negative
TN, FP, FN, TP = conf_matrix.ravel() # Extract values
```

8.4 Calculate sensitivity and specificity [4 marks]

```
# Calculate sensitivity
sensitivity = TP / (TP + FN) if (TP + FN) > 0 else 0

# Calculate specificity
specificity = TN / (TN + FP) if (TN + FP) > 0 else 0 # Specificity
```

8.5 Calculate precision and recall [4 marks]

```
# Calculate precision
precision = TP / (TP + FP) if (TP + FP) > 0 else 0 # Precision

# Calculate recall
recall = TP / (TP + FN) if (TP + FN) > 0 else 0 # Recall
```


Conclusion

The logistic regression model performs very well on the validation set, achieving high accuracy (93%), sensitivity (93%), specificity (93%), and precision (95%). These results indicate that the model effectively balances identification of positive and negative cases, with minimal misclassification.